

Business Problem - Predict the Price of Bangalore House

Using Linear Regression - Supervised Machine Learning Algorithm

Load Libraries

```
In [1]: import pandas as pd
```

Load Data

```
In [2]: # path = r"https://drive.google.com/uc?export=download&id=1xxDtrZKfuWQfL-6KA9XEd_eatitNPnKB"  
# df = pd.read_csv(path)
```

```
In [202]: df = pd.read_csv("bangalore house price prediction OHE-data.csv")
```

In [203]: df.head()

Out[203]:

	bath	balcony	price	total_sqft_int	bhk	price_per_sqft	area_typeSuper built-up Area	area_typeBuilt- up Area	area_typePlot Area	avail
0	3.0	2.0	150.0	1672.0	3	8971.291866	1	0	0	
1	3.0	3.0	149.0	1750.0	3	8514.285714	0	1	0	
2	3.0	2.0	150.0	1750.0	3	8571.428571	1	0	0	
3	2.0	2.0	40.0	1250.0	2	3200.000000	1	0	0	
4	2.0	2.0	83.0	1200.0	2	6916.666667	0	0	1	

In [163]: pd.set_option("display.max_columns",None)
pd.set_option("display.max_rows",None)

In [166]: df.head()

Out[166]:

	bath	balcony	price	total_sqft_int	bhk	price_per_sqft	area_typeSuper built-up Area	area_typeBuilt- up Area	area_typePlot Area	avail
0	3.0	2.0	150.0	1672.0	3	8971.291866	1	0	0	
1	3.0	3.0	149.0	1750.0	3	8514.285714	0	1	0	
2	3.0	2.0	150.0	1750.0	3	8571.428571	1	0	0	
3	2.0	2.0	40.0	1250.0	2	3200.000000	1	0	0	
4	2.0	2.0	83.0	1200.0	2	6916.666667	0	0	1	

Split Data

```
In [167]: X = df.drop('price', axis=1)
y = df['price']

print('Shape of X = ', X.shape)
print('Shape of y = ', y.shape)
```

```
Shape of X = (7120, 107)
Shape of y = (7120,)
```

```
In [168]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=51)

print('Shape of X_train = ', X_train.shape)
print('Shape of y_train = ', y_train.shape)
print('Shape of X_test = ', X_test.shape)
print('Shape of y_test = ', y_test.shape)
```

```
Shape of X_train = (5696, 107)
Shape of y_train = (5696,)
Shape of X_test = (1424, 107)
Shape of y_test = (1424,)
```

Feature Scaling

```
In [169]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
sc.fit(X_train)
X_train = sc.transform(X_train)
X_test = sc.transform(X_test)
```

Linear Regression - ML Model Training

```
In [170]: from sklearn.linear_model import LinearRegression
          lr = LinearRegression()

          lr.fit(X_train, y_train)
```

Out[170]: LinearRegression()

```
In [171]: lr.coef_
```

Out[171]: array([-5.70206143e+00, -1.25679916e+00, 8.27341833e+01, -1.44906911e+01,
 5.75662723e+01, 1.88468905e-01, -1.72593897e+00, -4.51058311e+00,
 -2.22589244e+00, -4.28978455e+00, -2.44590976e+00, 5.40246226e-01,
 -1.03633400e+00, 1.43064873e+00, -6.25029424e-02, -1.51548783e+00,
 -2.14422789e-01, 2.16244155e+00, -1.48710228e+00, 1.95250816e+00,
 -3.10761125e+00, -1.28138668e+00, -1.01367155e+00, 1.37968545e-02,
 1.10383858e+00, 1.26497611e+00, -3.52405517e+00, -1.21398741e+00,
 -5.04622019e-01, 1.46299181e+00, -5.50064233e-01, -8.46468162e-02,
 6.84882188e-01, -1.39849820e+00, -1.94761710e-02, -1.57716300e+00,
 4.20886278e-01, 8.03443207e-01, 2.99182164e+00, 3.86430413e-03,
 1.05037261e-01, 2.89115612e-01, -3.16916626e-01, 1.05625868e+00,
 -1.39649279e+00, -3.10533604e+00, 1.01764011e-01, -7.49672917e-02,
 -8.03271555e-01, -1.27061856e+00, -8.54046164e-01, 2.64566484e-01,
 9.10688839e-01, -8.23059458e-01, -9.07215234e-01, 1.22059216e+00,
 2.11418894e+00, -5.38187400e-01, -1.32164338e+00, -8.28349340e-01,
 1.28167980e+00, -1.92911295e-01, 6.65824485e-02, 3.65563139e-02,
 -1.85069853e+00, 1.49068024e+00, -9.57964753e-01, -9.36110163e-01,
 -7.45634897e-01, 7.22643165e-02, -6.79260144e-01, -1.70853833e-01,
 -1.72288643e+00, -1.15833746e+00, 5.78931788e-01, 1.37836966e+00,
 -1.14424496e+00, 3.96188294e-01, -6.08013157e-01, -2.20959218e+00,
 3.45270810e-01, 1.01747431e-03, 1.06563895e-01, 3.04728530e+00,
 2.09496392e+00, -8.13481923e-01, -4.18437282e-01, 2.30993396e+00,
 3.31858800e-02, 8.07865914e-02, 5.37064987e-02, 1.55347699e+00,
 8.13889657e-01, -1.14636462e+00, 3.41805788e-01, -8.28022037e-01,
 1.68897360e+00, 2.97657524e-01, 9.59437517e-01, 4.57297702e-01,
 -2.22729515e-01, -1.48290835e+00, -6.26342867e-01, 5.86538254e-01,
 -1.78547310e+00, 2.19020231e-01, -3.45032599e-01])

```
In [172]: lr.intercept_
```

```
Out[172]: 95.0802729985955
```

Predict the value of Home and Test

```
In [173]: X_test[0, :]
```

```
Out[173]: array([ 0.71301986,  0.0112734 ,  0.30202307,  0.65677518, -0.48064341,  
                -1.7385623 ,  2.11587407, -0.25430867,  0.51007548, -0.18373025,  
                -0.16389438, -0.1473229 , -0.13023539, -0.12812824, -0.12598816,  
                -0.12454231, -0.12953656, -0.12381344, -0.12010681, -0.11551113,  
                -0.10992018, -0.10909925, -0.10660036, -0.11234866, -0.09315135,  
                -0.08618799, -0.08923672, -0.09023078, -0.08721571, -0.09023078,  
                -0.08721571, -0.08195215, -0.08195215, -0.07633675, -0.0751646 ,  
                -0.08085949, -0.0739743 , -0.07975227, -0.07153563, -0.0751646 ,  
                -0.0677166 , -0.08085949, -0.07153563, -0.07862985, -0.0751646 ,  
                -0.07862985, -0.06504853, -0.0751646 , -0.06901264, -0.0751646 ,  
                -0.06901264, -0.07028523, -0.07276497, -0.07028523, -0.06367332,  
                -0.06226825, -0.06226825, -0.06639573, -0.06504853, -0.05935999,  
                -0.06083125, -0.06639573, -0.06639573, -0.06226825, -0.06367332,  
                -0.05935999, -0.06639573, -0.06367332, -0.06226825, -0.06226825,  
                -0.05935999, -0.05935999, -0.05935999, -0.05630391, -0.05935999,  
                -0.05785186, -0.05935999, -0.05935999, -0.06083125, -0.06083125,  
                -0.05471275, -0.06083125, -0.06226825, -0.05935999, -0.05935999,  
                -0.06226825, -0.06226825, -0.05785186, -0.06504853, -0.06226825,  
                -0.06083125, -0.05935999, -0.05307449, -0.05630391, -0.06226825,  
                -0.05471275, -0.05935999, -0.05471275, -0.05471275, -0.05138463,  
                -0.05307449, -0.05307449, -0.05471275, -0.05471275, -0.05630391,  
                -0.05630391, -0.05138463])
```

```
In [174]: lr.predict([X_test[0, :]])
```

```
Out[174]: array([76.90661876])
```

```
In [175]: lr.predict(X_test)
```

```
Out[175]: array([ 76.90661876,  15.25005377, 113.6828165 , ...,  21.30296864,  
                71.43462962, 230.0414626 ])
```

```
In [176]: y_test
```

```
Out[176]: 2435      80.000  
          3113      40.000  
          426     120.000  
          1124      79.000  
          1161      45.000  
          1724      39.000  
          1110      53.000  
          3591      78.000  
          1978     101.000  
          4383      62.000  
          4602     105.000  
          3536     246.000  
           820     198.000  
          4500     233.000  
          5550     145.000  
          1273      48.000  
          1905      86.000  
          5007      94.000  
          4447     160.000  
          6580      54.000
```

```
In [177]: lr.score(X_test, y_test)
```

```
Out[177]: 0.7903837092682249
```

Ridge and lasso Regression in python AI.ipynb

```
In [178]: from sklearn.linear_model import Ridge ,Lasso
```

```
In [179]: re=Ridge()
```

```
In [180]: re.fit(X_train,y_train)
```

```
Out[180]: Ridge()
```

```
In [181]: re.score(X_test,y_test)
```

```
Out[181]: 0.7905686374336595
```

```
In [182]: ls=Lasso()
```

```
In [183]: ls.fit(X_train,y_train)
```

```
Out[183]: Lasso()
```

```
In [184]: ls.score(X_test,y_test)
```

```
Out[184]: 0.8036373003525779
```

```
In [185]: re2=Ridge(alpha=2)  
re2.fit(X_train,y_train)  
re2.score(X_test,y_test)
```

```
Out[185]: 0.7907530260397906
```

```
In [186]: re9=Ridge(alpha=99)  
re9.fit(X_train,y_train)  
re9.score(X_test,y_test)
```

```
Out[186]: 0.8063391305505513
```

```
In [187]: re125=Ridge(alpha=199)
re125.fit(X_train,y_train)
re125.score(X_test,y_test)
```

```
Out[187]: 0.8184305724743911
```

```
In [188]: re225=Ridge(alpha=230)
re225.fit(X_train,y_train)
re225.score(X_test,y_test)
```

```
Out[188]: 0.8215074786386392
```

```
In [189]: ls2=Lasso(alpha=2)
ls2.fit(X_train,y_train)
ls2.score(X_test,y_test)
```

```
Out[189]: 0.8160181533703599
```

```
In [190]: ls3=Lasso(alpha=3)
ls3.fit(X_train,y_train)
ls3.score(X_test,y_test)
```

```
Out[190]: 0.8263450613017429
```

```
In [191]: re225.predict(X_test)
```

```
Out[191]: array([ 80.05325774,  17.34381545, 117.32762318, ...,  18.95418427,
        69.90269234, 224.78244508])
```

```
In [192]: y_test
```

...


```
In [209]: y_predict=re225.predict([X_test[3, :]])  
y_predict
```

```
Out[209]: array([83.19266762])
```

```
In [210]: y_test[1124]
```

```
Out[210]: 79.0
```

```
In [ ]:
```

Root mean square Error

```
In [213]: y_predict=re225.predict(X_test)  
y_predict
```

```
Out[213]: array([ 80.05325774,  17.34381545, 117.32762318, ...,  18.95418427,  
                  69.90269234, 224.78244508])
```

In [214]:

y_test

Out[214]:

2435	80.000
3113	40.000
426	120.000
1124	79.000
1161	45.000
1724	39.000
1110	53.000
3591	78.000
1978	101.000
4383	62.000
4602	105.000
3536	246.000
820	198.000
4500	233.000
5550	145.000
1273	48.000
1905	86.000
5007	94.000
4447	160.000
5500	54.000

In [215]:

```
from sklearn.metrics import mean_squared_error
import numpy as np
mse=mean_squared_error(y_test,y_predict)
rmse=np.sqrt(mse)
print("MSE=",mse)
print("RMSE=",rmse)
```

MSE= 3586.438971337255
RMSE= 59.88688480241108